

# ***Embedded Systems***

## ***Instruction Set Architecture Design***

### ***Lesson 03***

**Francesco Menichelli**

**francesco.menichelli@uniroma1.it**

# Instruction Set Architecture (ISA)

The instruction set architecture is concerned with:

- Organization of programmable storage (memory & registers):  
Includes the amount of addressable memory and number of available registers.
- Data Types & Data Structures: Encodings & representations.
- Instruction Set: What operations are specified.
- Instruction formats and encoding.
- Modes of addressing and accessing data items and instructions
- Exception conditions.

# Computer Instruction Sets

- Regardless of computer type, CPU structure, or hardware organization, every machine instruction must specify the following:
  - Opcode: Which operation to perform. Example: add, load, and branch.
  - Where to find the operand or operands, if any: Operands may be contained in CPU registers, main memory, or I/O ports.
  - Where to put the result, if there is a result: May be explicitly mentioned or implicit in the opcode.
  - Where to find the next instruction: Without any explicit branches, the instruction to execute is the next instruction in the sequence or a specified address in case of jump or branch instructions.

# Operation Types in The Instruction Set

Operator Type	Examples
<b>Arithmetic and logical</b>	Integer arithmetic and logical operations: add, or
<b>Data transfer</b>	Loads-stores (move on machines with memory addressing)
<b>Control</b>	Branch, jump, procedure call, and return, traps.
<b>System</b>	Operating system call/return, virtual memory management instructions ...
<b>Floating Point</b>	Floating point operations: add, multiply, divide
<b>Decimal</b>	Decimal add, decimal multiply, decimal to character conversion
<b>String</b>	String move, string compare, string search
<b>Media</b>	The same operation performed on multiple data (e.g. Intel MMX, SSE)

# RISC vs. CISC

- **Reduced instruction set computer (RISC):**
  - load/store;
    - operands in memory must be first loaded into register before any operation
  - 32-bit fixed format instruction
  - 32 32-bit General Purpose Registers
    - (R0 contains zero, DP take pair of Registers)
  - 3-address, reg-reg arithmetic instruction
  - Single address mode for load/store: base + displacement
    - No indirection
  - Simple branch conditions
  - Examples: ARM, MIPS, Sun SPARC, PowerPC
- **Complex instruction set computer (CISC):**
  - many addressing modes;
    - can directly operate on operands in memory
  - many operations.
  - variable instruction length
  - Examples: Intel x86, 68000, VAX microprocessors and compatibles

# Instruction Set Architecture: What Must be Specified?


- Instruction Format or Encoding
  - How is it decoded?
- Addressing of operands and result
  - Immediate addressing
  - Register addressing
  - Direct addressing
  - Register indirect addressing
  - Indexed addressing
  - Based-indexed addressing
  - PC Relative
  - Stack addressing
- Data type and Size
  - Fixed, Float, Vector
  - 8, 16, 32, 64 bit
  - Little Endian vs. Big Endian
- Operations
  - ADD, MUL, DIV, SHIFT, OR, AND
- Successor instruction
  - Jumps, conditions, branches
  - Fetch-decode-execute is implicit! –Von Neumann

# Instruction Format and Encoding

- **Considerations affecting instruction set encoding:**
  - The number of registers and addressing modes supported by ISA.
  - The impact of the size of the register and addressing mode fields on the average instruction size and on the average program.
  - To encode instructions into lengths that will be easy to handle in the implementation. On a minimum to be a multiple of bytes (8 bits).
- **Instruction Encoding Classification:**
  - Fixed length encoding: Faster and easiest to implement in hardware.
  - Variable length encoding: Produces smaller instructions.
  - Hybrid encoding.

# Three Examples of Instruction Set Encoding




## Variable Length Encoding: VAX (1-53 bytes)



## Fixed Length Encoding: MIPS, PowerPC, SPARC (all instructions are 4 bytes each)



## Hybrid Encoding: IBM 360/370, Intel 80x86





# Types of ISAs According To Operand Addressing Fields

- **Memory-To-Memory Machines:**

- Operands obtained from memory and results stored back in memory by any instruction that requires operands.
- No local CPU registers are used in the CPU datapath.
- Include:
  - The 4-Address ISA.
  - The 3-Address ISA.
  - The 2-Address ISA.

- **The 1-address (Accumulator) ISA:**

- A single local CPU special-purpose register (accumulator) is used as the source of one operand and as the result destination

# Types of ISAs According To Operand Addressing Fields

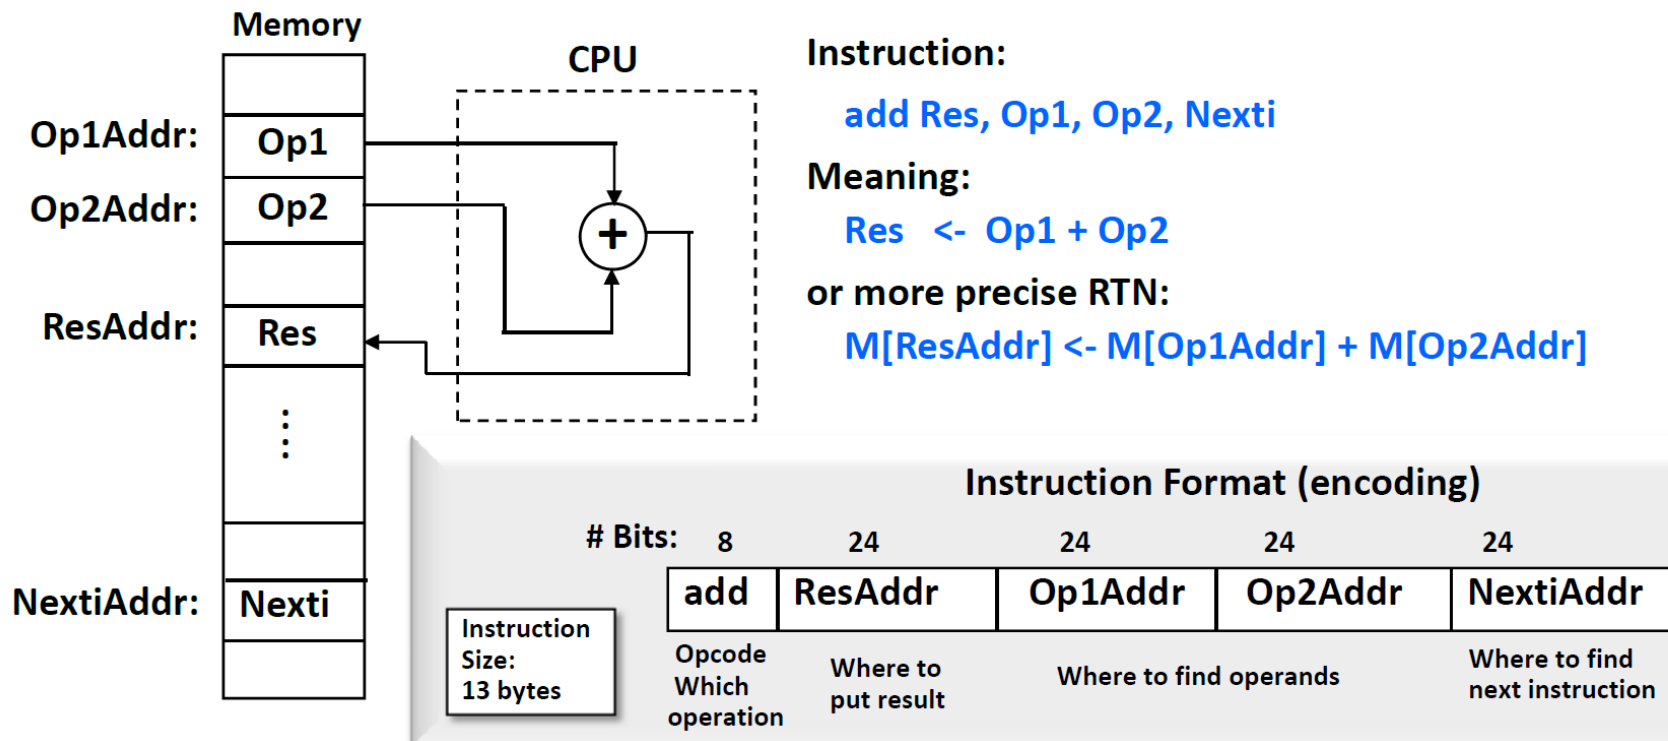
- **The 0-address or Stack ISA:**
  - A push-down stack is used in the CPU.
- **General Purpose Register (GPR) ISA:**
  - The CPU datapath contains several local general-purpose registers which can be used as operand sources and as result destinations.
  - A large number of possible addressing modes.
  - Load-Store or Register-To-Register Machines: GPR machines where only data movement instructions (loads, stores) can obtain operands from memory and store results to memory.

# Types of ISAs -Memory-To-Memory Machines: The 4-Address ISA

No program counter (PC) or other CPU registers are used.

Instruction encoding has four address fields to specify:

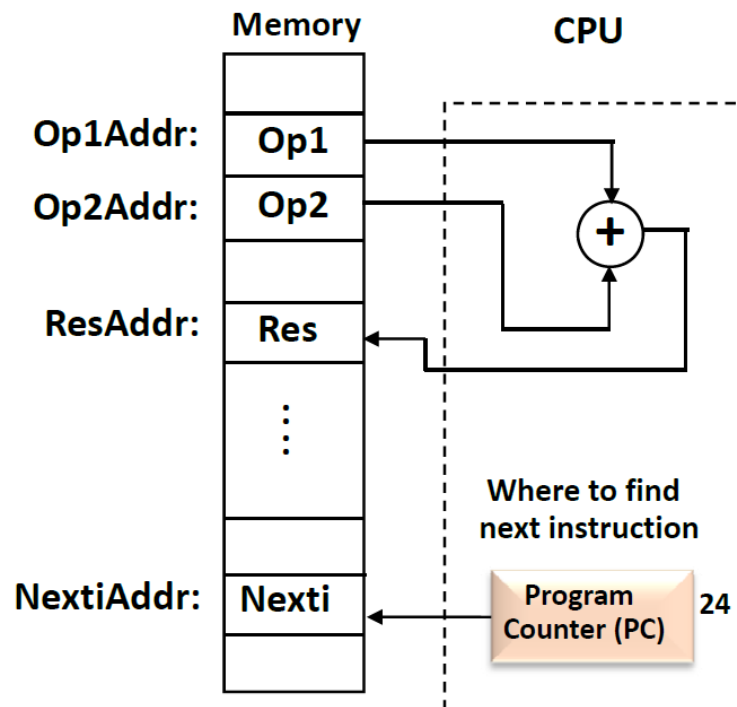
- Location of first operand, location of second operand.
- Place to store the result, location of next instruction.



# Types of ISAs -Memory-To-Memory Machines: The 3-Address ISA

A program counter (PC) is included within the CPU which points to the next instruction.

No CPU storage (general-purpose registers).



Instruction:

**sub Res, Op1, Op2**

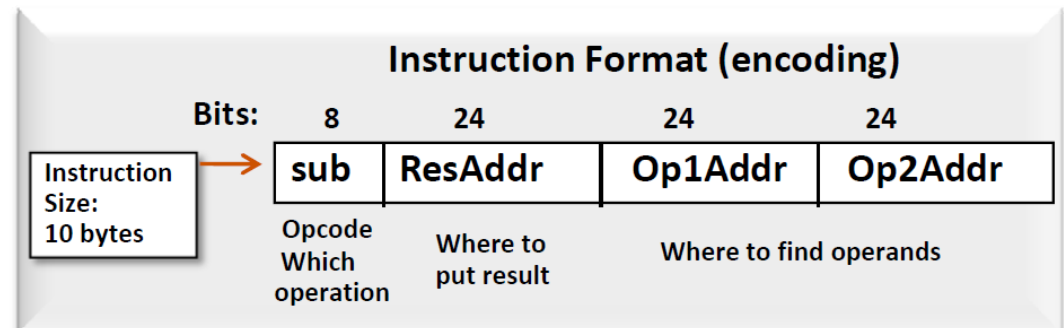
Meaning:

**Res  $\leftarrow$  Op1 - Op2**

or more precise RTN:

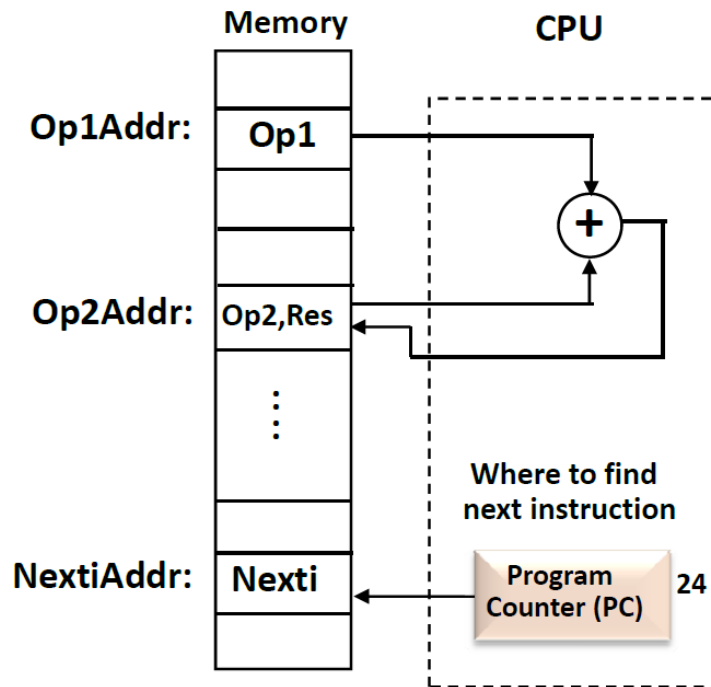
**$M[\text{ResAddr}] \leftarrow M[\text{Op1Addr}] - M[\text{Op2Addr}]$**

**PC  $\leftarrow$  PC + 10** Increment PC



# Types of ISAs -Memory-To-Memory Machines: The 2-Address ISA

The 2-address Machine: Result is stored in the memory address of one of the operands.



Can address  $2^{24}$  bytes = 16 MBytes

Instruction:

**add Op2, Op1**

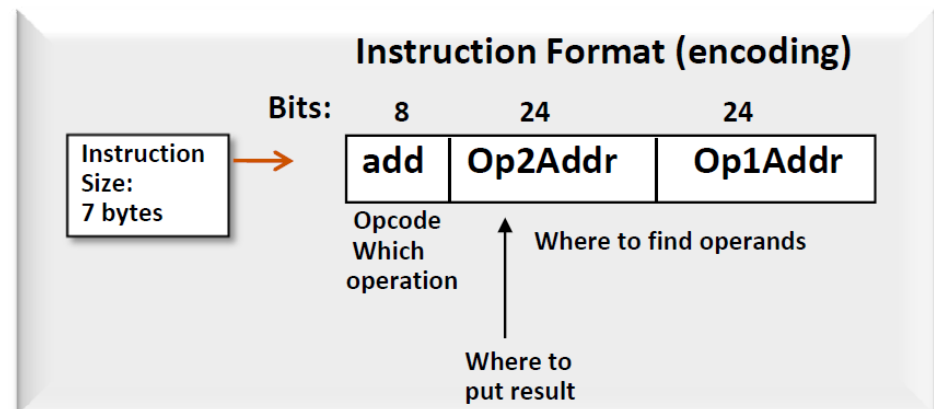
Meaning:

**Op2 <- Op1 + Op2**

or more precise RTN:

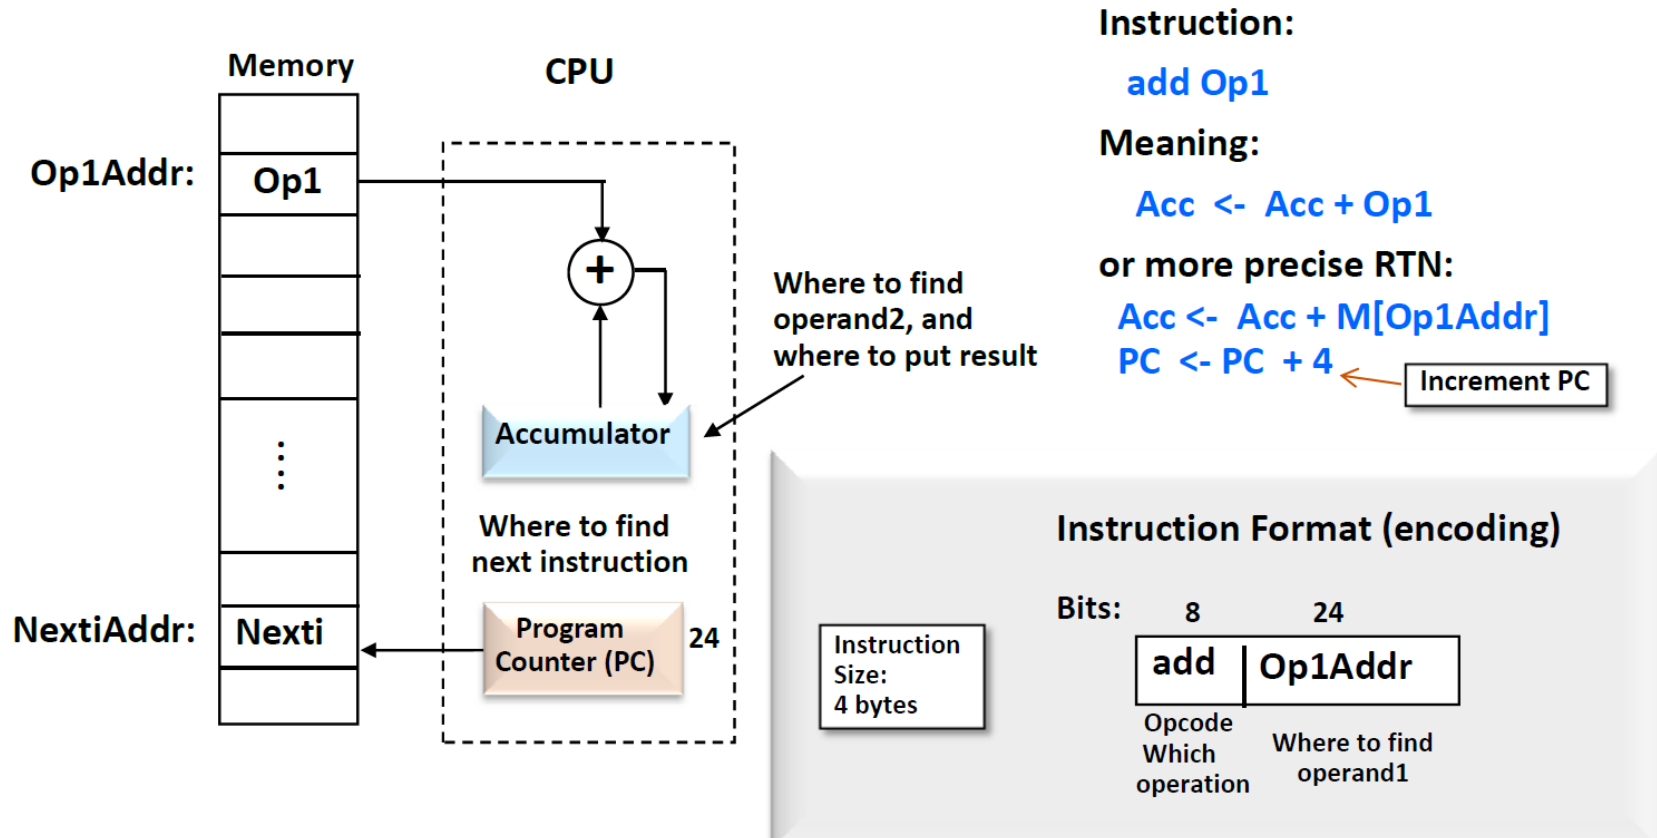
**$M[Op2Addr] \leftarrow M[Op1Addr] + M[Op2Addr]$**

**$PC \leftarrow PC + 7$**  ← **Increment PC**



# Types of ISAs -Memory-To-Memory Machines: The 1-Address(Accumulator) ISA

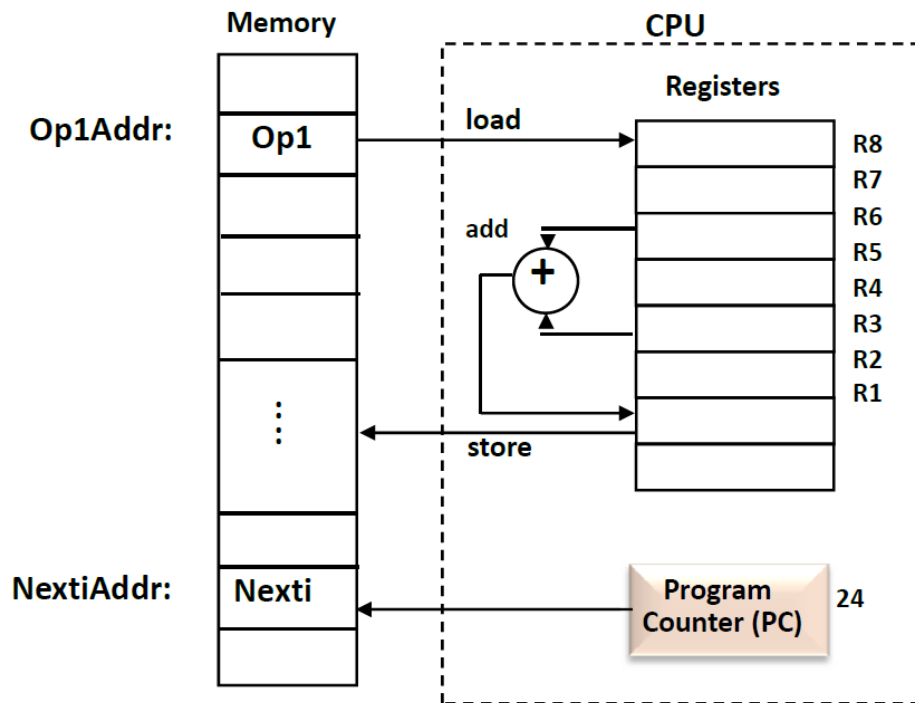
A single accumulator in the CPU is used as the source of one operand and result destination.



# Types of ISAs General Purpose Register (GPR) Machines

CPU contains several general-purpose registers which can be used as operand sources and result destination.

Eight general purpose Registers (GPRs) assumed here: R1-R8



The add instruction has three register specifier fields  
While load, store instructions have one register specifier field and one memory address specifier field

Instruction:

**load R8, Op1**

Meaning:

**$R8 \leftarrow M[Op1Addr]$**

**$PC \leftarrow PC + 5$**

Instruction:

**add R2, R4, R6**

Meaning:

**$R2 \leftarrow R4 + R6$**

**$PC \leftarrow PC + 3$**

Instruction:

**store R2, Op2**


Meaning:

**$M[Op2Addr] \leftarrow R2$**

**$PC \leftarrow PC + 5$**

Instruction Format

Bits: 8 3 24




Opcode

Where to find operand1

Size = 4.375 bytes rounded up to 5 bytes

Instruction Format

Bits: 8 3 3 3




Opcode

Des Operands

Size = 2.125 bytes rounded up to 3 bytes

Instruction Format

Bits: 8 3 24



Opcode

Destination

Size = 4.375 bytes rounded up to 5 bytes

# Expression Evaluation Example with 3-, 2-, 1-Address, and GPR Machines

For the expression  $A = (B + C) * D - E$  where A-E are in memory

3-Address	2-Address	1-Address Accumulator	GPR Register-Memory	GPR Load-Store
add A, B, C mul A, A, D sub A, A, E	load A, B add A, C mul A, D sub A, E	load B add C mul D sub E store A	load R1, B add R1, C mul R1, D sub R1, E store A, R1	load R1, B load R2, C add R3, R1, R2 load R1, D mul R3, R3, R1 load R1, E sub R3, R3, R1 store A, R3
3 instructions Code size: 30 bytes 9 memory accesses for data	4 instructions Code size: 28 bytes 11 memory accesses for data	5 instructions Code size: 20 bytes 5 memory accesses for data	5 instructions Code size: 25 bytes 5 memory accesses for data	8 instructions Code size: 34 bytes 5 memory accesses for data



# Instruction Set Architecture Tradeoffs

- 3-address machine: shortest code sequence; a large number of bits per instruction; large number of memory accesses.
- General purpose register machine (GPR):
  - Addressing modified by specifying among a small set of registers with using a short register address (all new ISAs since 1975).
  - Advantages of GPR:
    - Low number of memory accesses. Faster, since register access is currently still much faster than memory access.
    - Registers are easier for compilers to use.
    - Shorter, simpler instructions.
- Load-Store Machines: GPR machines where memory addresses are only included in data movement instructions (loads/stores) between memory and registers.

# Instruction Set Architecture Tradeoffs

- **Accumulator Machines**
  - Requires storing lots of temporary and intermediate values in memory
  - Accumulator is only really beneficial for a chain (sequence) of calculations where the result of one is the input to the next.
- **Memory-to-Memory ISAs**
  - Main memory much slower than arithmetic circuits
    - This was as true in 1950 as in 2015!
  - It takes a lot of room to specify memory addresses
  - Results are often used one or two instructions later

# Typical GPR ISA Memory Addressing Modes

Addressing mode	Sample Instruction	Meaning
Register	add R4, R3	$R4 \leftarrow R4 + R3$
Immediate	add R4, #3	$R4 \leftarrow R4 + 3$
Displacement	add R4, 10 (R1)	$R4 \leftarrow R4 + \text{Mem}[10 + R1]$
Indirect	add R4, (R1)	$R4 \leftarrow R4 + \text{Mem}[R1]$
Indexed	add R3, (R1 + R2)	$R3 \leftarrow R3 + \text{Mem}[R1 + R2]$
Absolute	add R1, (1001)	$R1 \leftarrow R1 + \text{Mem}[1001]$
Memory indirect	add R1, @ (R3)	$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R3]]$
Autoincrement	add R1, (R2) +	$R1 \leftarrow R1 + \text{Mem}[R2]$ $R2 \leftarrow R2 + d$
Autodecrement	add R1, -(R2)	$R2 \leftarrow R2 - d$ $R1 \leftarrow R1 + \text{Mem}[R2]$
Scaled	add R1, 100 (R2) [R3]	$R1 \leftarrow R1 + \text{Mem}[100 + R2 + R3 * d]$

“d” is generally the size of the memory word width e.g. 4 bytes

# Addressing Modes

- **Register addressing:**

- Specify register number rather than address.

MOV R1, R2

- Copy content of register 2 to register 1.

- **Immediate addressing:**

- Address part of operand contains operand itself.

MOV R1, #4 -> Load constant 4 to register 1.

- Only small integer constants can be specified in this way.
- No memory reference to fetch data
- Fast
- Limited range



# Addressing Modes

- **Register indirect addressing:**
  - Operand address is not contained in instruction but in a register.
  - Operand address is a pointer.
    - ADD R1, (R2); add to register R1 word at address contained in R2.
  - Can refer to different addresses in different instruction.
    - Example: assembly code for adding the elements of an array.


```
MOV R1, #0      ; accumulate sum in R1, initially 0
MOV R2, #A      ; R2 = address of the array A
MOV R3, #A+4096 ; R3 = address of first word beyond A
LOOP: ADD R1, (R2) ; register indirect through R2 to get operand
      ADD R2, #4   ; increment R2 by one word (4 bytes)
      CMP R2, R3   ; are we done yet?
      BLT LOOP     ; if R2 < R3, we are not done, so continue
```

# Addressing Modes

- **Indexed addressing:**

- Memory is addressed by giving a register plus a constant offset.
- Example: processing of static arrays.

MOV R4, A(R2); load into R1 word whose address has offset A from content of R2.



Array is at a fixed address; register contains current index.

- Example: assembly code for computing:  $\sum_i A_i * B_i$ .

```
MOV R1, #0      ; accumulate the sum in R1, initially 0
MOV R2, #0      ; R2 = index i
MOV R3, #4096   ; R3 = first index value not in use
LOOP: MOV R4, A(R2) ; R4 = A[i]
      MUL R4, B(R2) ; R4 = A[i] * B[i]
      ADD R1, R4    ; sum all the products into R1
      ADD R2, #4    ; i = i+4 (1 word = 4 bytes)
      CMP R2, R3    ; are we done yet?
      BLT LOOP     ; if R2 < R3, we are not done, so continue
```

# Addressing Modes

- **Based-Indexed Addressing**

- Address is computed by sum of two registers plus optional offset.
- Processing of dynamic arrays.
  - MOV R4, (R2+R5); load into R4 word whose address is the sum of R2 and R5.
- R5 is the base address of the array.
- R2 is the current index.
- Replace loop code in previous example as follows:

```
...  
    MOV R5, #A      ; R5 = address of A  
    MOV R6, #B      ; R6 = address of B  
LOOP: MOV R4, (R2+R5) ; R4 = A[i]  
      MUL R4, (R2+R6) ; R4 = A[i] * B[i]  
...  

```

# Addressing Modes

- **Memory Indirect**

- Memory cell pointed to by address field contains the address of (pointer to) the operand  
EA = (A) ; Look in A, find address (A) and look there for operand  
e.g. ADD (A)  
Add contents of cell pointed to by contents of A to accumulator
- May be nested, multilevel, cascaded  
e.g. EA = (((A)))
- Multiple memory accesses to find operand
- Very slow



# Addressing Modes for Branch Instructions

- **How to specify target address of branch instructions/procedure calls?**
  - Direct addressing: unconditional branches (gotos).
    - Generated from conditionals and loops.
  - Register indirect addressing or indexed mode.
    - Program may compute target address (computed goto, switch).
  - PC-relative addressing: indexed mode where PC acts as register.
    - Target address is specified as offset to current instruction.
- **Modes presented so far are also useful for branch instructions**

# Considerations in designing an ISA

- **Architectural State: Memory and Registers**
  - Abstraction used by both compiler and Microarchitecture
  - Register specialization?
  - RISC and VLIW tend to use a large pool of general-purpose registers
  - Implementation techniques provide illusion of large number of registers or fast flat memory space (register renaming)
- **Pipelining and operational latency**
  - Sequential execution vs. pipelined execution model
  - Exposed or hidden pipelining (latency of operations): delayed branches, bypass networks, scoreboard
  - DSPs enforce uniform latency model
- **Delayed branches expose part of the RISC pipeline**

# Considerations in designing an ISA

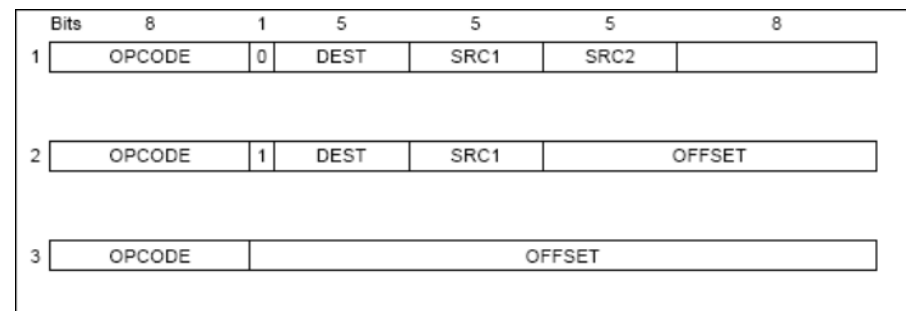
- **Encoding and Architectural Style:**
  - Architecture description contains architectural state, instruction set, and execution model
    - Implementations are free to build whatever structure they want, as long as they simulate the architectural model
  - RISC Encodings
    - Single instruction corresponds to a single operation
    - 1:1 encoding (Instruction encoding size: Number of run-time operations)
  - CISC Encodings
    - RISC subset plus more complex instructions
    - Variable: Dynamic encoding
  - VLIW Encodings
    - n:n; Fallacy: VLIW Instructions Have Fixed Width
  - DSP Encodings
    - 1:n
  - Vector Encodings 1:variable

# Considerations in designing an ISA

- In a clean design, every opcode should permit every addressing mode.

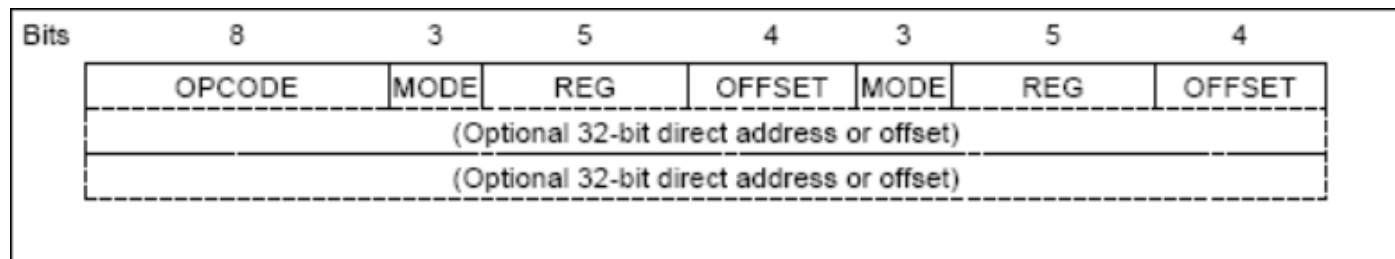
- Three Address Machine

- Two formats selected by bit.
- 1 special format for branches



- Two Address Machine

- Each operand specified by 12 bits.
- Mode, register, offset.
- Optional 32-bit word for address.



# Summary

- **Programmable CPU cores are important for the control parts of typical embedded application.**
- **They are well supported with tools to support the development of end-user software. (vs. deeply embedded SW)**
- **“Keep it Simple” heuristic (RISC vs. CISC)**
  - Make frequent cases fast and rare cases correct.
  - Regular (orthogonal) instruction set
  - No special features that match a high level language construct.
  - At least 16 registers to ease register allocation.
- **Embedded cores are often light cores which are a compromise between performance, area and power dissipation. (vs. stand-alone CPU cores which are optimized for performance)**

# CPU in embedded systems

- Main constraints in embedded systems
  - Cost (Area)
  - Performance
  - Power consumption
- They are not independent
  - Performance/Cost
  - Performance/Power consumption

# CPU in embedded systems

- Very wide range of applications
  - disposable electronic, consumer, automotive, Safety, Space
- Can not be achieved by a single or a reduced set of architectures
  - 4 bit architectures (very low end)
  - 8bit – 32bit (the most common)
  - 16bit, DSP, VLIW (specialized applications)
- The same constraints impact the whole architecture
  - Bus width
  - Memory subsystem (flash, ram, cache)
  - Power management

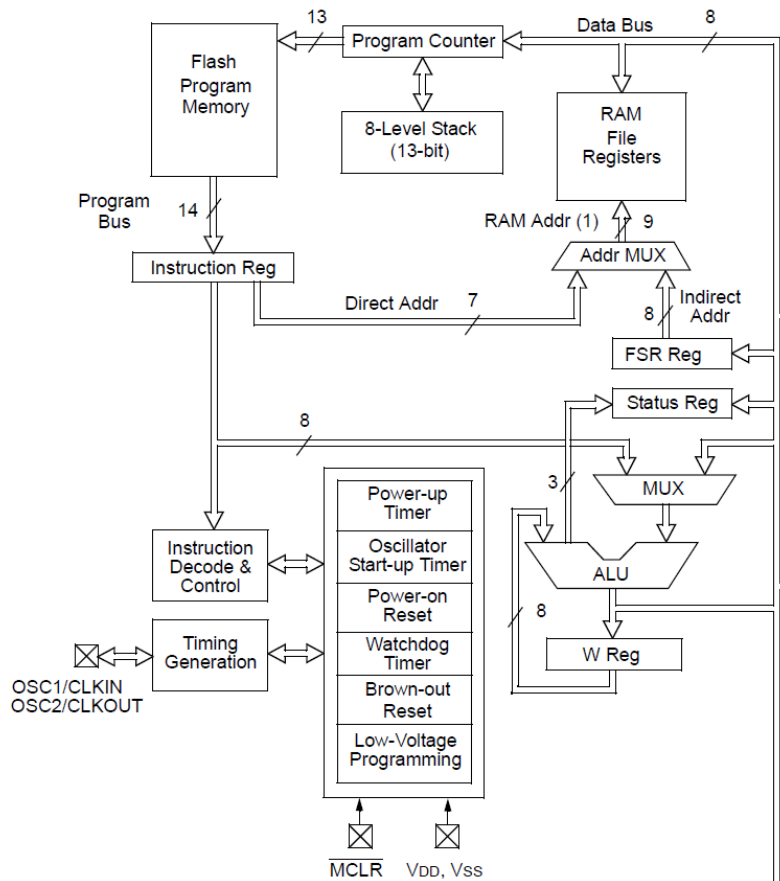
# CPU in embedded systems

- The result is a multitude of components and architectures
  - Microchip: PIC16, PIC18, PIC24, dsPIC, PIC32
  - ST Microelectronics: **ST6, ST7, ST9, ST10**, STM8, STM32
- The same can be repeated for Atmel, TI (Texas Instruments), Infineon, NXP, etc.



# Microchip PIC16F architecture

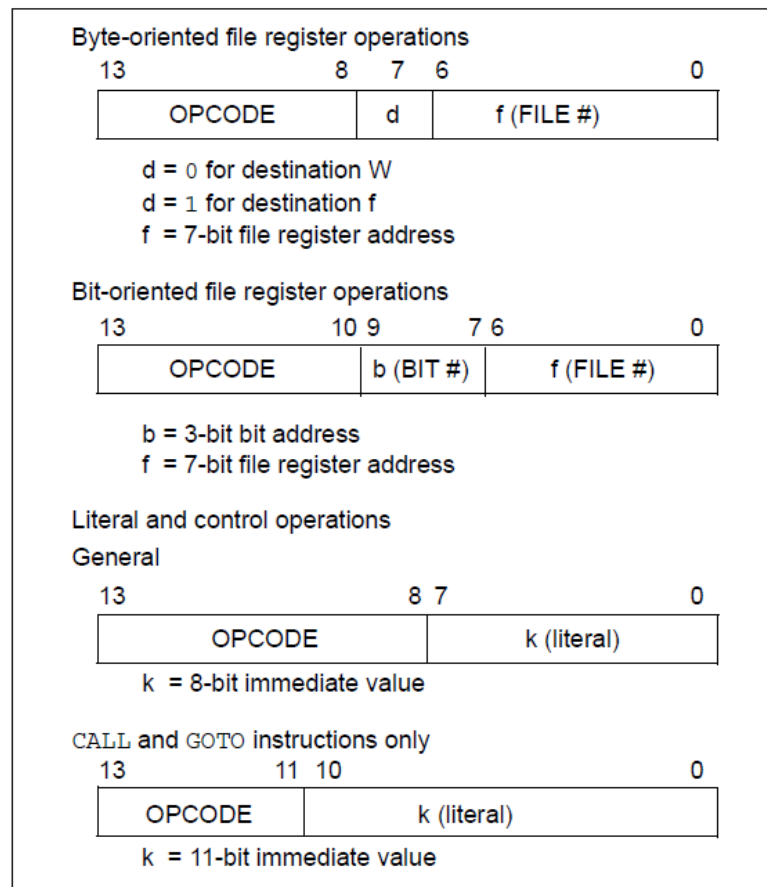
- Very low cost, in-house design (no compatibility with other CPUs)



- Strict Harvard architecture
- “File registers” RAM (64-256byte)
- Two stage pipeline (fetch, execute)
- 14-bit instruction set (14-bit instruction memory)
- 9-bit memory addressing (banked)
- 8-bit integer ALU
- Single data register (W)
- Very limited Stack (8-Level, PC only)

# Microchip PIC16F instruction format

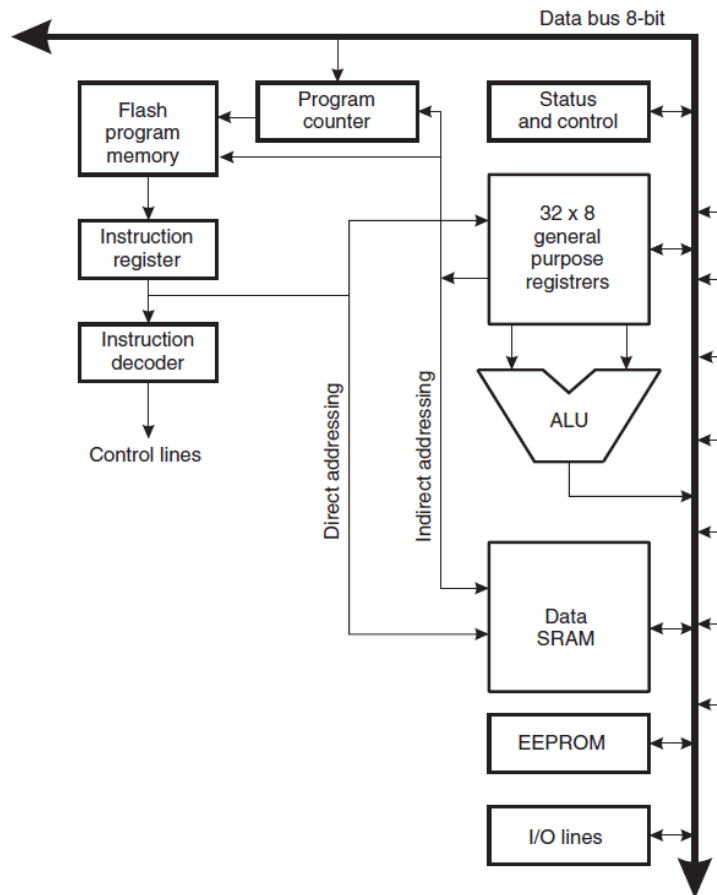
**FIGURE 15-1: GENERAL FORMAT FOR INSTRUCTIONS**



- 14 bit instruction set
- Accumulator based architecture (1-address ISA)
- Destination selectable (accumulator, memory)
- 4 instructions formats
  - byte oriented (MOV, ADD, AND, etc.)
  - bit oriented (set/reset bits)
  - literal (immediate accumulator loading)
  - Control flow (call, jump)

# ATMEL AVR8 architecture

- Low cost, in-house design (no compatibility with other CPUs)



- Harvard architecture
- data SRAM (1-4Kbyte)
- Two stage pipeline (fetch, execute)
- 16-bit instruction set (16-bit instruction memory)
- 16-bit memory addressing (linear)
- 8-bit integer ALU (two independent operands)
- 32 data registers (8-bit)
- RAM Stack (stack pointer)

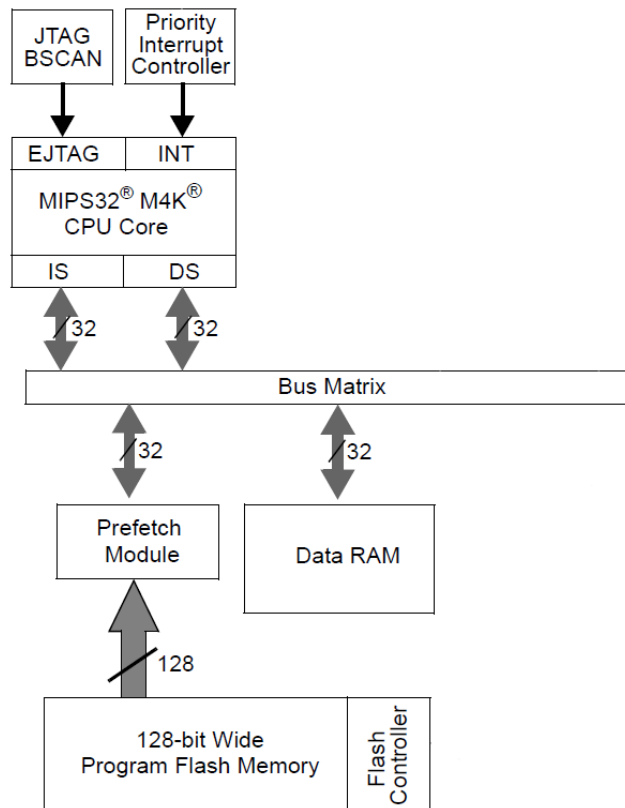
# ATMEL AVR8 architecture

- AVR CPU general purpose working registers

[illegible]

# Microchip PIC32 architecture

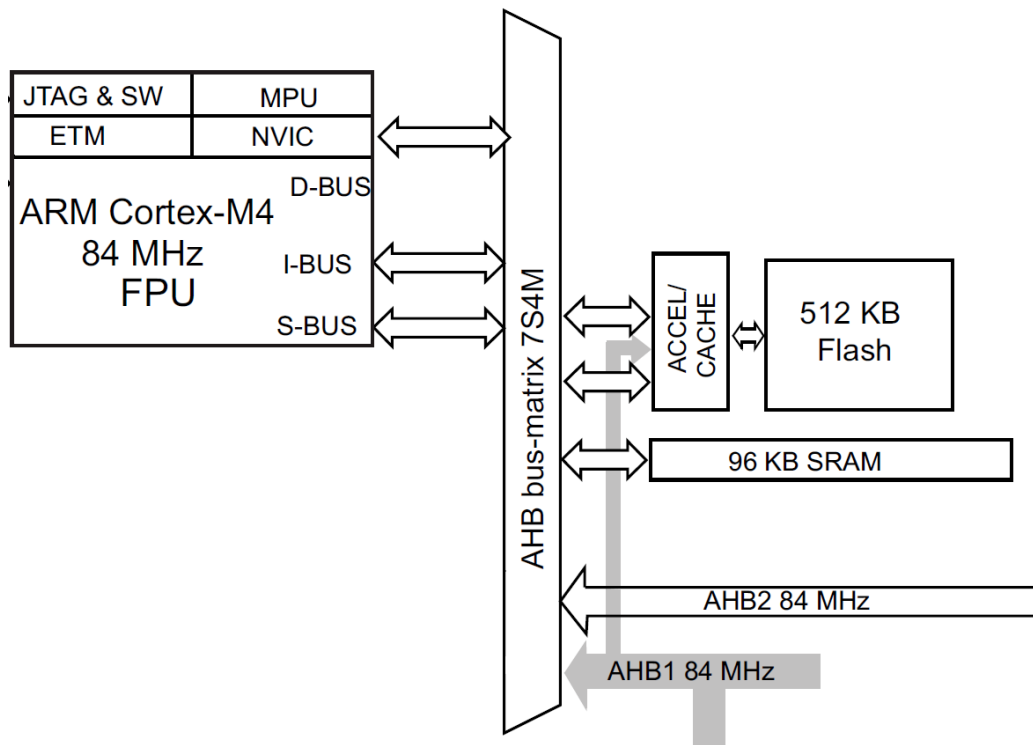
- MIPS32 M4K (MIPS Technologies)



- Harvard architecture (cache level)
- data SRAM (128-256Kbyte)
- five stage pipeline
- 32-bit instruction set (32-bit wide data memory)
- 128-bit wide program memory+prefetch
- 32-bit memory addressing (linear)
- 32-bit integer ALU (two independent operands+result)
- 32 data registers (32-bit)
- RAM Stack (stack pointer)

# STM32 architecture

- ARM Cortex-M4 (ARM Holdings plc): ARMv7E-M architecture



- Harvard architecture (without cache)
- Separated D-BUS, I-BUS, S-BUS
- data SRAM
- five stage pipeline
- 16/32-bit instruction set (THUMB-2 instruction set)
- 32-bit memory addressing (linear)
- 32-bit integer ALU (two independent operands+result)
- 16 data registers (32-bit)
- RAM Stack (stack pointer)